

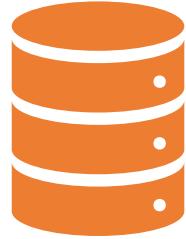
INF1001: Introduction to Computing

L2b: Database –

SQL



Topic outline: Database-SQL



Relational Queries & SQL

Overview

ER-Schema-SQL



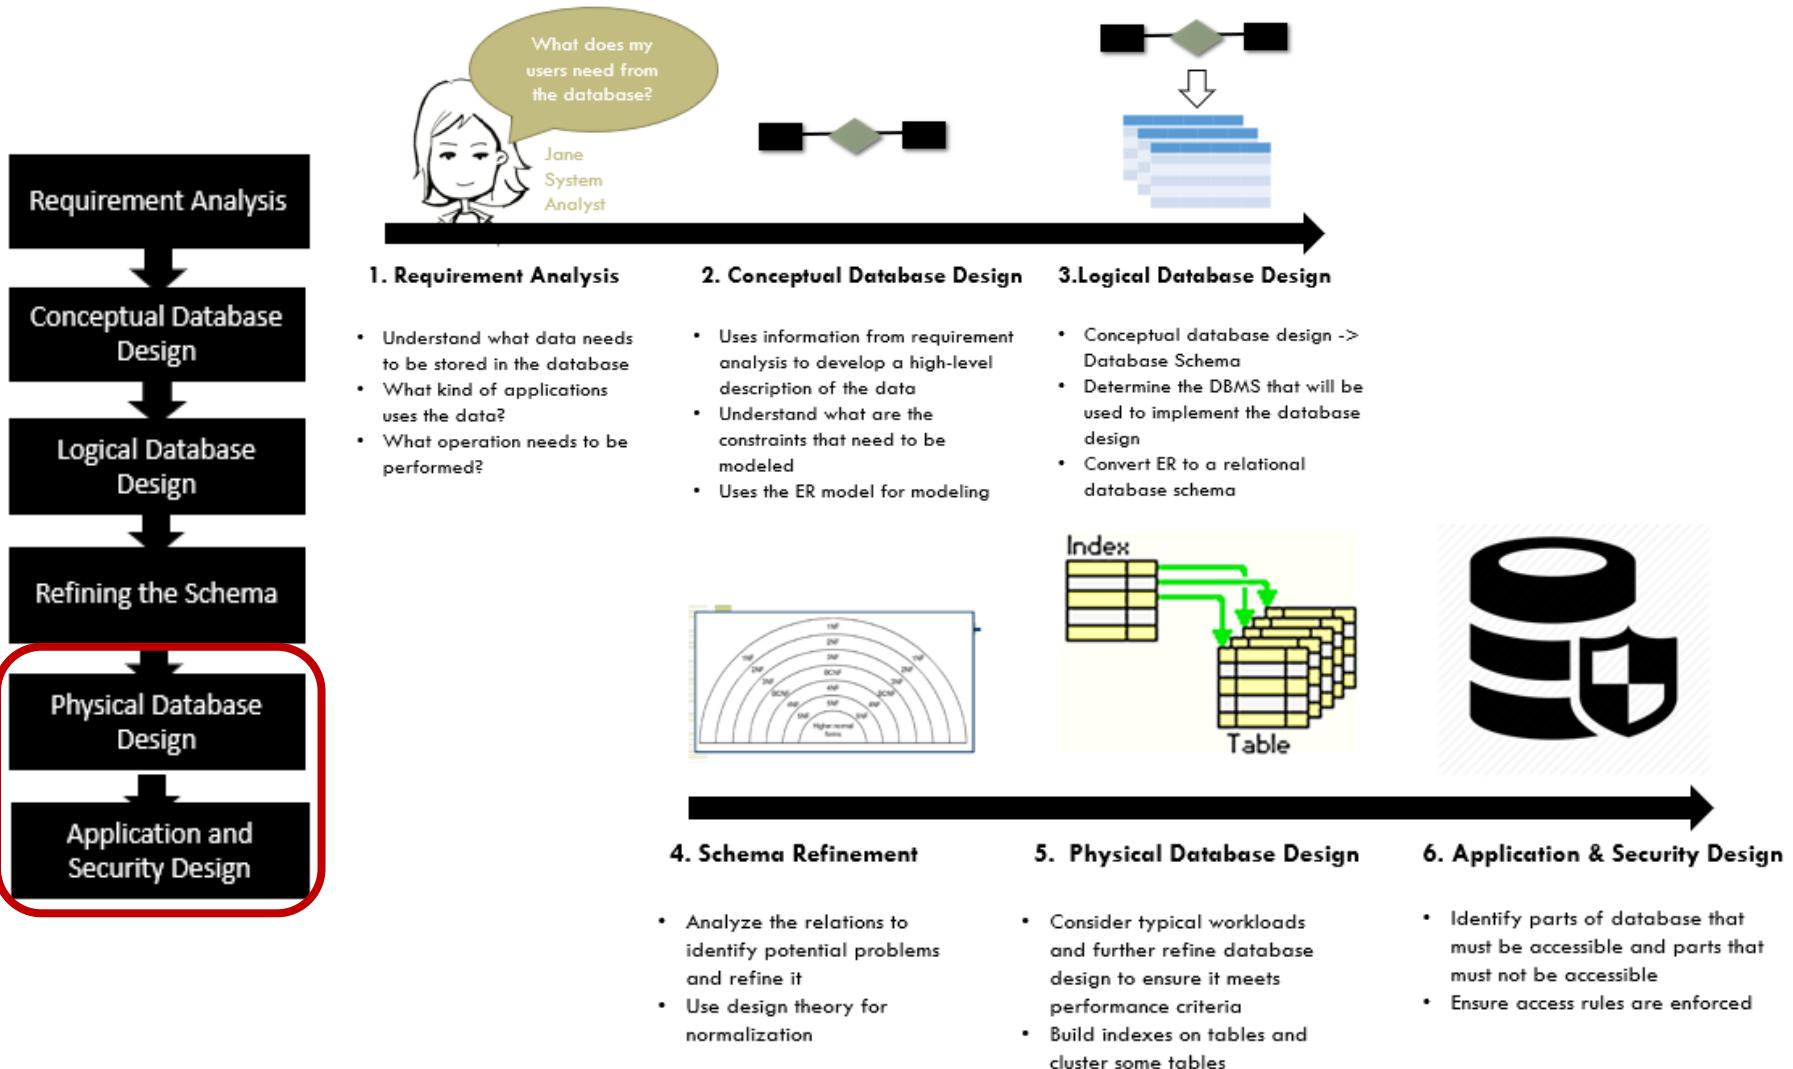
SQL Statements

**CREATE, INSERT, SELECT, Nested
Queries**

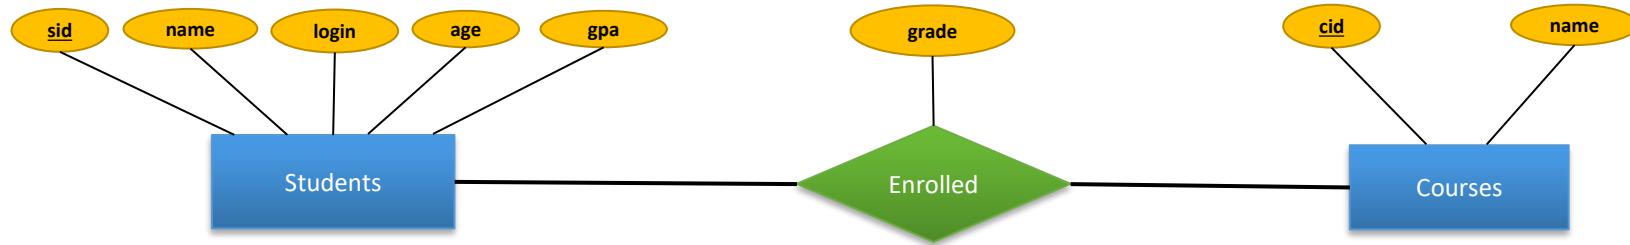
Syntax

Examples

Recap on Stages of Database Design



ER Diagram and Relational Schema



Relational Schema:

Students(sid varchar(20), name varchar(20), login varchar(10), age integer, gpa real)

Courses(cid varchar(20), name varchar(20))

Enrolled (cid varchar(20), sid varchar(20), grade varchar(5))

Relational Queries

- A major strength of the relational model: supports simple, powerful ***querying*** of data.
- Queries can be written intuitively, and the DBMS is responsible for ***efficient*** evaluation.
- Query:
 - Input and output of a query is a relation
 - Evaluated using instances of each input relation
 - Produces ***instance*** of output relation

Structured Query Language (SQL)

- Most widely used commercial relational database (DB) language
- Originally developed by IBM
- SQL is a very-high-level language.
 - Say “**what** to do” rather than “**how** to do it.”
 - Avoid a lot of data-manipulation details needed in procedural languages like C++ or Java.
- Database management system figures out “best” way to execute query.
 - Called “query optimization.”

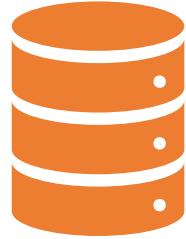


Structured Query Language (SQL)

- SQL statements:

- CREATE
- INSERT
- SELECT
- Nested Queries

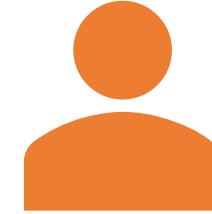
Outline



Relational Queries & SQL

Overview

ER-Schema-SQL



SQL Statements

CREATE, INSERT, SELECT, Nested Queries

Syntax

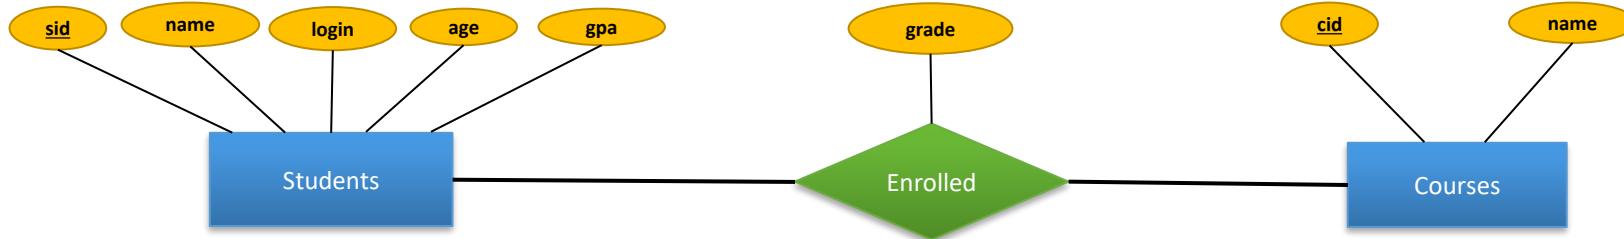
Examples

1. SQL - CREATE

- Creating tables/relations

```
CREATE TABLE table-name (
    attribute-name1 attribute type1,
    ...,
    attribute-namen attribute typen
    PRIMARY KEY (attribute-name)
)
```

1. SQL – CREATE



Schema

```
Students(sid varchar(20),
         name varchar(20),
         login varchar(10),
         age integer,
         gpa real)
```

```
CREATE TABLE Students
  (sid VARCHAR(20),
   name VARCHAR(20),
   login VARCHAR(10),
   age INTEGER,
   gpa REAL)
```

- Creates the Students relation.
- Observe that the type (domain) of each field is specified, and enforced by the DBMS whenever tuples are added or modified.

1. SQL – CREATE

```
CREATE TABLE Students
  (sid VARCHAR(20),
   name VARCHAR(20),
   login VARCHAR(10),
   age INTEGER,
   gpa REAL,
   PRIMARY KEY (sid))
```

Primary Key:

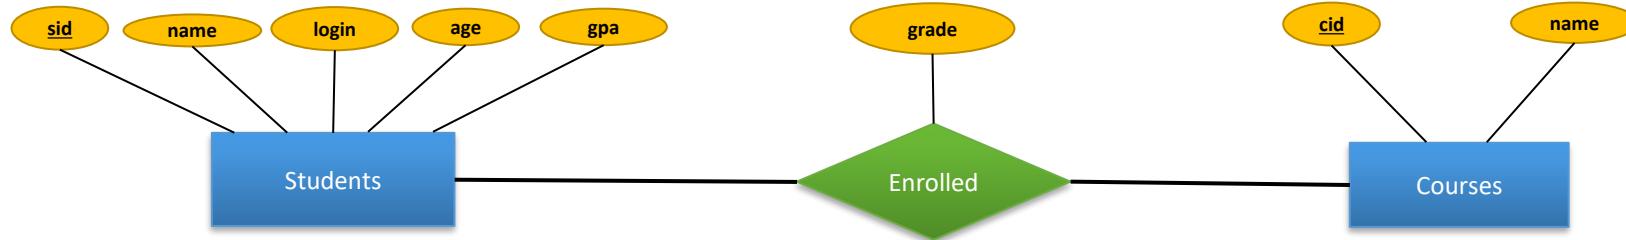
- Uniquely identifies each record in a table
- Must contain unique values
- Cannot contain NULL values

Table can have **only 1 primary key**, which may consist of a single or multiple fields

Students

| <u>sid</u> | name | login | age | gpa |
|------------|-------|------------|-----|-----|
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53688 | Smith | smith@eecs | 18 | 3.2 |
| 53650 | Smith | smith@math | 19 | 3.8 |

1. SQL – CREATE



Schema

```
Courses(cid varchar(20), name varchar(20))
```

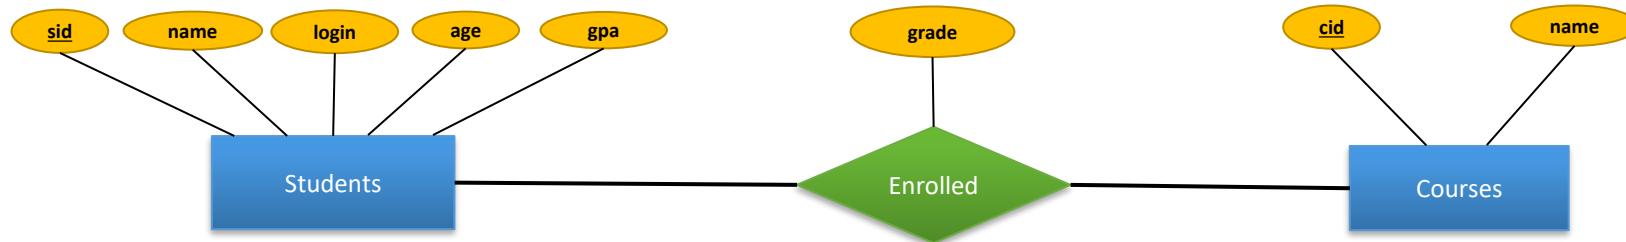
```
CREATE TABLE Courses
  (cid VARCHAR(20),
   name VARCHAR(20),
   PRIMARY KEY (cid))
```

Courses

| <u>cid</u> | name |
|------------|-----------------------|
| ICT1001 | Introduction to ICT |
| ICT1002 | Programming |
| ICT1003 | Computer Architecture |

1. SQL – CREATE

- Enrolled table holds information about courses that students take.



`Students(sid varchar(20), name varchar(20), login varchar(10), age integer, gpa real)`

`Courses(cid varchar(20), name varchar(20))`

`Enrolled (sid varchar(20), cid varchar(20), grade varchar(5))`

```

CREATE TABLE Students
  (sid VARCHAR(20), name VARCHAR(20),
   login VARCHAR(10), age INTEGER, gpa REAL,
   PRIMARY KEY (sid) )
  
```

```

CREATE TABLE Courses
  (cid VARCHAR(20),
   name VARCHAR(20),
   PRIMARY KEY (cid) )
  
```

```

CREATE TABLE Enrolled
  (sid VARCHAR(20), cid VARCHAR(20),
   grade VARCHAR(5))
  
```

Only students listed in the Students relation should be allowed to enroll for courses. How to enforce this?

1. SQL – CREATE: Foreign Key

- A FOREIGN KEY is a key used to link two tables together.
- A FOREIGN KEY is a field (or collection of fields) in one table that refers to the PRIMARY KEY in another table.
 - The table with the foreign key is called the **referencing** or **child** table
 - Table with the primary key is called the **referenced** or **parent** table.
- The FOREIGN KEY constraint is used to prevent actions (such as deleting tuples) that would destroy links between tables.
- The FOREIGN KEY constraint also prevents invalid data from being inserted into the foreign key column, because it has to be one of the values contained in the parent table that it points to.
 - E.g., FAIL/ERROR if trying to add a grade for the student **sid** in the child **Enrolled** table that references a non-existent **sid** in the parent **Students** table

1. SQL – CREATE: Foreign Key

- Only students listed in the Students relation should be allowed to enroll for courses.

```
CREATE TABLE Students
  (sid VARCHAR(20), name VARCHAR(20),
   login VARCHAR(10), age INTEGER, gpa REAL,
   PRIMARY KEY (sid))
```

```
CREATE TABLE Courses
  (cid VARCHAR(20),
   name VARCHAR(20),
   PRIMARY KEY (cid))
```

```
CREATE TABLE Enrolled
  (sid VARCHAR(20), cid VARCHAR(20), grade VARCHAR(5),
   PRIMARY KEY (sid, cid),
   FOREIGN KEY (sid) REFERENCES Students(sid)
   FOREIGN KEY (cid) REFERENCES Courses(cid))
```

Enrolled is the referencing/child relation, while Students is the referenced/parent relation

Enrolled

| <u>sid</u> | <u>cid</u> | grade |
|------------|------------|-------|
| 53666 | ICT1001 | C |
| 53666 | ICT1002 | B |
| 53650 | ICT1001 | A |
| 53666 | ICT1003 | B |

| <u>sid</u> | name | login | age | gpa |
|------------|-------|------------|-----|-----|
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53688 | Smith | smith@eecs | 18 | 3.2 |
| 53650 | Smith | smith@math | 19 | 3.8 |

1. SQL – CREATE: Foreign Key

- If all foreign key constraints are enforced, **referential integrity** is achieved
- Consider parent **Students** and child **Enrolled** tables; ***sid*** in Enrolled is a foreign key that references Students relation.
 - What should be done if an **Enrolled** tuple with a non-existent student ***sid*** is **inserted**?
 - Reject it!
 - What should be done if a **Students** tuple is **deleted**? Several options:
 - Disallow deletion of the Students tuple that is referred to
 - Delete all Enrolled tuples that refer to it
 - Set ***sid*** in Enrolled tuples that refer to it to a default ***sid*** (for e.g. 00000)
 - Similar actions are taken if primary key Student is **updated**

Logical Database Design (Schema)

Requirement Analysis

Conceptual Database
Design

Logical Database
Design

Refining the Schema

Physical Database
Design

Application and
Security Design

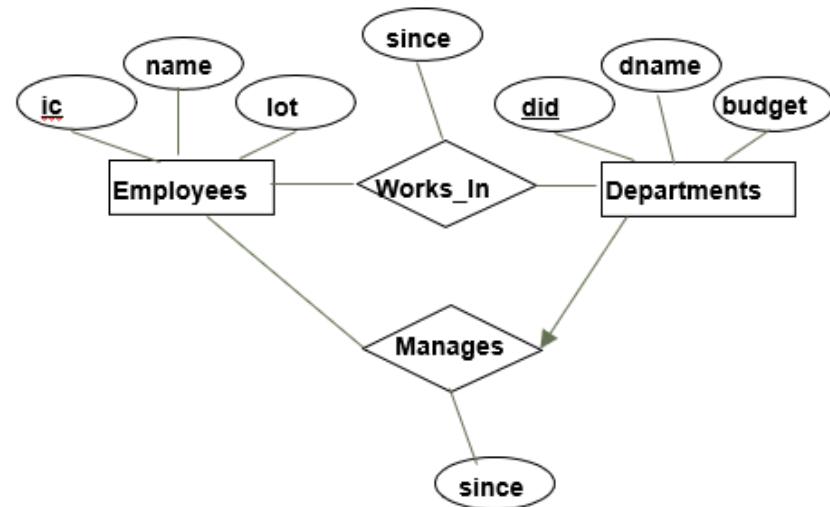
Employee Management System

We would like to develop a database application to keep information of our Employees.

We would like to keep track of the employee information such IC, name, and parking lot.

We would also like to keep track of departments through their unique ID, name and budget. In addition, we would like to keep track of the department that the employee is working in and the time the employee joined.

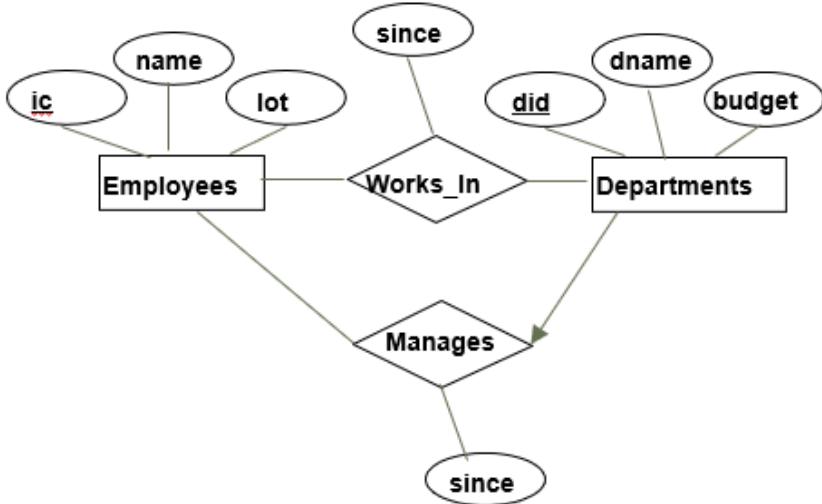
A department is managed by one employee. The same employee can manage multiple departments.



Relational Schema

Employees (ic: char(11), name: char(11), lot: integer)
Works_In (ic: char(11), did: int, since: date)
Dept_Mgr (did: int, dname: char(11), budget: real,
 ic: char(11), since: date)

From Logical Database Design (Schema) to SQL



```

CREATE TABLE Employees (
    ic VARCHAR(11),
    name VARCHAR(11),
    lot INTEGER,
    PRIMARY KEY (ic)
)
  
```

```

CREATE TABLE Dept_Mgr (
    did INTEGER,
    dname VARCHAR(11),
    budget REAL,
    ic VARCHAR(11),
    since DATE,
    PRIMARY KEY (did),
    FOREIGN KEY (ic) REFERENCES Employees(ic)
)
  
```

Relational Schema

```

Employees (ic: char(11), name: char(11), lot: integer)
Works_In (ic: char(11), did: int, since: date)
Dept_Mgr (did: int, dname: char(11), budget: real,
          ic: char(11), since: date)
  
```

```

CREATE TABLE Works_In (
    ic VARCHAR(11),
    did INTEGER,
    since DATE,
    PRIMARY KEY (ic, did),
    FOREIGN KEY (ic) REFERENCES Employees(ic),
    FOREIGN KEY (did) REFERENCES Dept_Mgr(did)
)
  
```

Example Instances

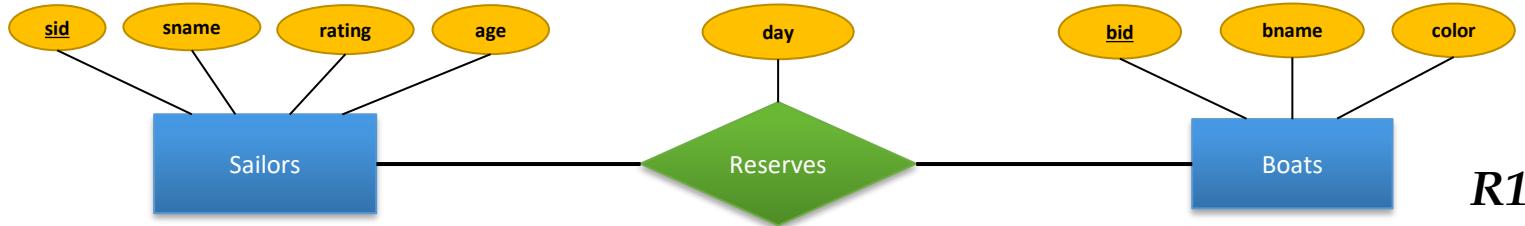
- We will use these instances of the Sailors, Boats and Reserves relations in our examples.

```

Sailors(sid: integer, sname: string, rating: integer, age: real)

Boats(bid: integer, bname: string, color: string)

Reserves(sid: integer, bid: integer, day: date)
  
```



```

CREATE TABLE Sailors (sid int unsigned, sname varchar(10), rating int, age real,
PRIMARY KEY (sid))
  
```

Sailors

| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 22 | Dustin | 7 | 45.0 |
| 31 | Lubber | 8 | 55.5 |
| 58 | Rusty | 10 | 35.0 |
| 74 | Rusty | 7 | 35.0 |

Example Instances

```
CREATE TABLE Boats (bid int unsigned, bname varchar(10), color varchar(10),
PRIMARY KEY (bid))
```

Boats

| <u>bid</u> | <u>bname</u> | <u>color</u> |
|------------|--------------|--------------|
| 101 | Interlake | Blue |
| 102 | Interlake | Red |
| 103 | Clipper | green |
| 104 | Marine | red |

```
Sailors(sid: integer, sname: string, rating: integer, age: real)

Boats(bid: integer, bname: string, color: string)

Reserves(sid: integer, bid: integer, day: date)
```

```
CREATE TABLE Reserves (sid int unsigned, bid int unsigned, day date,
PRIMARY KEY(sid, bid),
FOREIGN KEY(sid) REFERENCES sailors(sid),
FOREIGN KEY(bid) REFERENCES boats(bid))
```

Reserves

| <u>sid</u> | <u>bid</u> | <u>day</u> |
|------------|------------|------------|
| 22 | 101 | 18/10/21 |
| 58 | 103 | 25/12/21 |

2. SQL - INSERT

- Adding tuples to relation

```
INSERT INTO table-name
    (attribute-name1, ..., attribute-namen,)
VALUES
    (attribute-value1, ..., attribute-valuen,)
```

- Example: insert a single tuple into Students relation:

```
INSERT INTO Students (sid, name, login, age, gpa)
VALUES (53688, 'Smith', 'smith@eeecs', 18, 3.2)
```

Students

| <u>sid</u> | name | login | age | gpa |
|------------|-------|-------------|-----|-----|
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53650 | Smith | smith@math | 19 | 3.8 |
| 53688 | Smith | smith@eeecs | 18 | 3.2 |

3. SQL - SELECT

```
SELECT column1, column2, ...
FROM table_name;
```

- The SELECT statement is used to select data from a database.
- The data returned is stored in a result table, called the *result-set*.
- *column1, column2, ...* are the field names of the table you want to select data from.

If you want to select all the fields available in the table:

```
SELECT * FROM table_name;
```

If you want to return only distinct (different) values:

```
SELECT DISTINCT column1, column2, ...
FROM table_name;
```

SQL SELECT Statement

```
CREATE TABLE Sailors (sid int unsigned, sname varchar(10), rating int, age real,  
PRIMARY KEY (sid))
```

Q1: Find the name and age of all sailors

```
SELECT sname, age  
FROM Sailors
```

Query Result:

| sname | age |
|--------|------|
| Dustin | 45.0 |
| Lubber | 55.5 |
| Rusty | 35.0 |
| Rusty | 35.0 |

Q2: Find DISTINCT name and age of all sailors

```
SELECT DISTINCT sname, age  
FROM Sailors
```

Query Result:

| sname | age |
|--------|------|
| Dustin | 45.0 |
| Lubber | 55.5 |
| Rusty | 35.0 |

Removes
duplicates in
the result set

SQL SELECT Statement

```
CREATE TABLE Sailors (sid int unsigned, sname varchar(10), rating int, age real,  
PRIMARY KEY (sid))
```

Q3: Find sid, name, age, rating of all sailors

```
SELECT sid, sname, age, rating  
FROM Sailors
```

OR

```
SELECT *  
FROM Sailors
```

Query Result:

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | Dustin | 7 | 45.0 |
| 31 | Lubber | 8 | 55.5 |
| 58 | Rusty | 10 | 35.0 |
| 74 | Rusty | 7 | 35.0 |

3. SQL - SELECT

- Adding conditions in your SELECT statement

```
SELECT desired attributes
FROM one or more tables
WHERE condition about tuples of tables
```

SQL SELECT Statement

Q4: Find all sailors with rating > 7

```
SELECT sname, age
FROM Sailors
WHERE rating > 7
```

Query Result:

| sname | age |
|--------|------|
| Lubber | 55.5 |
| Rusty | 35.0 |

Sailors

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | Dustin | 7 | 45.0 |
| 31 | Lubber | 8 | 55.5 |
| 58 | Rusty | 10 | 35.0 |
| 74 | Rusty | 7 | 35.0 |

Q5: Find all sailors who are older than 35

```
SELECT sname, age
FROM Sailors
WHERE age > 35
```

Query Result:

| sname | age |
|--------|------|
| Dustin | 45.0 |
| Lubber | 55.5 |

Sailors

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | Dustin | 7 | 45.0 |
| 31 | Lubber | 8 | 55.5 |
| 58 | Rusty | 10 | 35.0 |
| 74 | Rusty | 7 | 35.0 |

3. SQL – SELECT: Complex conditions

- Complex condition in WHERE clause
 - Boolean operators AND, OR, NOT.
 - Comparisons =, <>, <, >, <=, >=.
 - And many other operators that produce boolean-valued results.

3. SQL – SELECT: Complex conditions

- SELECT FROM 2 or more tables – Join Tables

INF1001 Aug 2023

Q6: Find the names of sailors who have reserved at least one boat

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid
```

JOIN-ing TWO tables

Sailors

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | Dustin | 7 | 45.0 |
| 31 | Lubber | 8 | 55.5 |
| 58 | Rusty | 10 | 35.0 |

Cross product

Reserves

| sid | bid | day |
|-----|-----|----------|
| 22 | 101 | 18/10/21 |
| 58 | 103 | 25/12/21 |

Evaluate the WHERE condition
 $S.Sid = R.sid$

| (sid) | sname | rating | age | (sid) | bid | day |
|-------|--------|--------|------|-------|-----|----------|
| 22 | Dustin | 7 | 45 | 22 | 101 | 18/10/21 |
| 22 | Dustin | 7 | 45 | 58 | 103 | 25/12/21 |
| 31 | Lubber | 8 | 55.5 | 22 | 101 | 18/10/21 |
| 31 | Lubber | 8 | 55.5 | 58 | 103 | 25/12/21 |
| 58 | Rusty | 10 | 35.0 | 22 | 101 | 18/10/21 |
| 58 | Rusty | 10 | 35.0 | 58 | 103 | 25/12/21 |

| sname |
|--------|
| Dustin |
| Rusty |

3. SQL – SELECT: Complex conditions

- SELECT FROM 2 or more tables – Join Tables

Summary : How is the query computed to produce a JOIN?

1. Produce a temporary table that contains the **cross product** of the two tables:
 - Each row from table 1 (**size m**) is concatenated with each row from table 2 (**size n**)
 - Produce a **size $m \times n$** cross product table
2. Eliminate rows from the cross product temporary table by evaluating the WHERE condition
3. Outputs final results using the attributes in the SELECT statement

3. SQL – SELECT: Complex conditions

- SELECT FROM 2 or more tables – Join Tables

Q7: Find the names of sailors who have reserved boat 103

```
SELECT S.sname
FROM Sailor S, Reserves R
WHERE S.sid=R.sid AND R.bid=103
```

JOIN-ing TWO tables

Sailors

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | Dustin | 7 | 45.0 |
| 31 | Lubber | 8 | 55.5 |
| 58 | Rusty | 10 | 35.0 |

Cross product

Reserves

| sid | bid | day |
|-----|-----|----------|
| 22 | 101 | 18/10/21 |
| 58 | 103 | 25/12/21 |

| (sid) | sname | rating | age | (sid) | bid | day |
|-------|--------|--------|------|-------|-----|----------|
| 22 | Dustin | 7 | 45 | 22 | 101 | 18/10/21 |
| 22 | Dustin | 7 | 45 | 58 | 103 | 25/12/21 |
| 31 | Lubber | 8 | 55.5 | 22 | 101 | 18/10/21 |
| 31 | Lubber | 8 | 55.5 | 58 | 103 | 25/12/21 |
| 58 | Rusty | 10 | 35.0 | 22 | 101 | 18/10/21 |
| 58 | Rusty | 10 | 35.0 | 58 | 103 | 25/12/21 |

| sname |
|-------|
| Rusty |

Evaluate the
WHERE condition
 $S.Sid = R.sid$

Evaluate the
WHERE condition
 $R.bid = 103$

3. SQL – SELECT: Complex conditions

- SELECT FROM 2 or more tables – Join Tables

Q8: Find the names of student who has obtained an A grade in their course

```
SELECT S.name, E.cid
FROM Students S, Enrolled E
WHERE S.sid=E.sid AND E.grade='A'
```

Students

| <u>sid</u> | <u>name</u> | <u>login</u> | <u>age</u> | <u>gpa</u> |
|------------|-------------|--------------|------------|------------|
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53688 | Smith | smith@eecs | 18 | 3.2 |
| 53650 | Smith | smith@math | 19 | 3.8 |

Enrolled

| <u>sid</u> | <u>cid</u> | <u>grade</u> |
|------------|-------------|--------------|
| 53666 | Carnatic101 | C |
| 53666 | Reggae203 | B |
| 53650 | Topology112 | A |
| 53666 | History105 | B |

We get:

| <u>S.name</u> | <u>E.cid</u> |
|---------------|--------------|
| Smith | Topology112 |

4. SQL – SELECT: Nested Queries

- A very powerful feature of SQL: a WHERE clause can itself contain an SQL query!
- To understand semantics of nested queries, think of a *nested loops* evaluation

4. SQL – SELECT: Nested Queries

Q9: Find the names of sailors who have reserved boat 103

Previous Example:
 (Q7) SELECT S.sname
 FROM Sailor S, Reserves R
 WHERE S.sid=R.sid AND R.bid=103

SELECT S.sname
 FROM Sailors S
 WHERE S.sid IN (SELECT R.sid
 FROM Reserves R
 WHERE R.bid=103)

IN tests whether a value is in a given set

Reserves

| Sid | bid | day |
|-----|-----|----------|
| 22 | 101 | 18/10/21 |
| 58 | 103 | 25/12/21 |

Sailors

| sid | sname | rating | age |
|-----|--------|--------|------|
| 22 | Dustin | 7 | 45.0 |
| 31 | Lubber | 8 | 55.5 |
| 58 | Rusty | 10 | 35.0 |

The nested subquery computes the multi set of sids for sailors who have reserved boat 103
 The top level query retrieves the names of sailors whose sid is in the set

- To find sailors who have not reserved #103, use NOT IN.
- To understand semantics of nested queries, think of a nested loops evaluation

More Examples

Q10: Find the sid of sailors who have reserved a **red** boat

```
SELECT R.sid
FROM Reserves R, Boats B
WHERE R.bid = B.bid AND B.color = 'red'
```

OR

```
SELECT R.sid
FROM Reserves R
WHERE R.bid IN (SELECT B.bid
                 FROM Boats B
                 WHERE B.color = 'red')
```

Sailors

| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 22 | Dustin | 7 | 45.0 |
| 31 | Lubber | 8 | 55.5 |
| 58 | Rusty | 10 | 35.0 |
| 74 | Rusty | 7 | 35.0 |

Boats

| <u>bid</u> | bname | color |
|------------|-----------|-------|
| 101 | Interlake | Blue |
| 102 | Interlake | Red |
| 103 | Clipper | green |
| 104 | Marine | red |

Reserves

| <u>sid</u> | <u>bid</u> | day |
|------------|------------|----------|
| 22 | 101 | 18/10/21 |
| 58 | 103 | 25/12/21 |

More Examples

Q11: Find the **names** of sailors who have reserved a **red boat**

```
SELECT S.sname
FROM   Sailors S, Reserves R, Boats B
WHERE  S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'
```

Sailors

| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 22 | Dustin | 7 | 45.0 |
| 31 | Lubber | 8 | 55.5 |
| 58 | Rusty | 10 | 35.0 |
| 74 | Rusty | 7 | 35.0 |

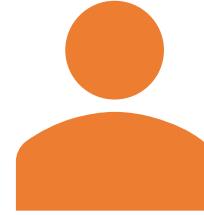
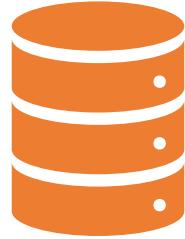
Boats

| <u>bid</u> | bname | color |
|------------|-----------|-------|
| 101 | Interlake | Blue |
| 102 | Interlake | Red |
| 103 | Clipper | green |
| 104 | Marine | red |

Reserves

| <u>sid</u> | <u>bid</u> | day |
|------------|------------|----------|
| 22 | 101 | 18/10/21 |
| 58 | 103 | 25/12/21 |

Summary



Relational Queries & SQL

Overview

ER-Schema-SQL

SQL Statements

**CREATE, INSERT, SELECT, Nested
Queries**

Syntax

Examples

To be continued in Year 2 Database courses (for CS and SE students)

More SQL practice

- Kueri.me
- <http://www.geektime.com/2016/07/24/kueri-development-like-google-translate-from-natural-language-to-sql/>
- <http://sqlzoo.net/>
- If you want answers, you can just add “?answer=1” at the end of the URL and the answers will display in the text boxes.
- <https://www.w3resource.com/sql-exercises/>

References

- Good reference: SQL Tutorial at

<https://www.w3schools.com/sql/>

w3schools.com

